

Frank-Wolfe with Variance-Controlled Step Size for Sparse Stochastic Optimization

Jiewen Luo Vaishnavi Nasery Rufina Lourdes

COMPSCI 651 – Optimization, University of Massachusetts Amherst

1 Introduction and Motivation

Stochastic Gradient Descent (SGD) and its adaptive variants are the dominant optimization tools in large-scale machine learning. However, in constrained settings where sparse solutions are required, they face two compounding difficulties. First, stochastic gradients carry high variance that destabilizes convergence, especially under a fixed learning rate [10]. Second, enforcing an ℓ_1 constraint through Euclidean projection requires solving an isotonic soft-thresholding subproblem at every iteration, costing $\mathcal{O}(d \log d)$ per step [2].

The Frank-Wolfe (FW) algorithm [4, 6] addresses the projection bottleneck by replacing it with a *linear minimization oracle* (LMO): at each step, it computes the direction that minimizes a linear approximation of the objective over the feasible set. For the ℓ_1 ball, this reduces to selecting the coordinate with the largest absolute gradient magnitude — an $\mathcal{O}(d)$ operation that automatically produces a sparse update direction. The resulting iterate remains feasible by construction via a convex combination update, entirely avoiding projection.

However, applying FW naively in the stochastic setting introduces a new challenge: the fixed or polynomial step size schedules from the deterministic theory are poorly matched to the high-variance, non-stationary gradient dynamics typical of mini-batch training. Overly large steps cause oscillation near the constraint boundary; overly small steps waste the sparsity benefit of the LMO by making negligible progress.

We propose **FW-VC**, a Frank-Wolfe variant with a *variance-controlled adaptive step size*. Rather than fixing the schedule a priori, FW-VC maintains an exponential moving average (EMA) of the squared gradient norm and scales the step size inversely to its square root — taking cautious steps when gradients are noisy and more aggressive steps when they stabilize. This design is motivated by the AdaGrad/RMSProp line of work [3, 12] adapted to the convex combination structure of FW updates.

Our central question mirrors the project proposal: *can a variance-controlled Frank-Wolfe method*

achieve stable convergence while producing meaningfully sparse solutions?

2 Related Work

The Frank-Wolfe algorithm, originally introduced for quadratic programming [4], was revisited by Jaggi [6], who showed that it achieves an $\mathcal{O}(1/T)$ convergence rate for smooth convex objectives and that the LMO solution for the ℓ_1 ball naturally induces sparsity. Hazan and Kale [5] extended the analysis to the online and stochastic settings, establishing $\mathcal{O}(1/\sqrt{T})$ convergence under bounded gradient noise.

Projection-free methods have since been applied to structured prediction [6], deep learning [9], and constrained recovery. Pokutta et al. [9] demonstrated FW updates in neural network training on ℓ_1 -constrained weight spaces, directly motivating our logistic regression setting.

Adaptive step-size methods for SGD — including AdaGrad [3], RMSProp [12], and Adam [7] — have shown strong empirical performance by scaling gradients by their historical magnitudes. Our work applies this idea to the FW framework, where the step size controls the convex interpolation coefficient rather than an additive gradient update. The compatibility requirements differ from the SGD case, as we discuss in Section 5.

Sparse optimization via ℓ_1 regularization (LASSO [11]) and ℓ_1 -constrained projection [2] provide the baselines against which we compare. The Robbins-Monro framework [10] provides the theoretical foundation for convergence conditions on the step size sequence.

3 Problem Formulation

We study the constrained stochastic optimization problem

$$\min_{x \in \mathbb{R}^d} f(x) = \mathbb{E}_\xi[f(x, \xi)] \quad \text{subject to} \quad \|x\|_1 \leq R, \quad (1)$$

where $f(\cdot, \xi)$ is a smooth convex loss (binary cross-entropy in our experiments), ξ is a random data point, and the ℓ_1 constraint promotes sparsity. In the logistic regression setting,

$$f(x, \xi) = -[y \log \sigma(a^\top x) + (1 - y) \log(1 - \sigma(a^\top x))],$$

where (a, y) is a feature-label pair, and $\sigma(z) = (1 + e^{-z})^{-1}$. The stochastic gradient on a mini-batch $\mathcal{B}$ of size B is

$$g_t = \frac{1}{B} X_{\mathcal{B}}^\top (\sigma(X_{\mathcal{B}} x_t) - y_{\mathcal{B}}).$$

4 Proposed Method: FW-VC

4.1 Frank-Wolfe LMO for the ℓ_1 Ball

At each iteration, the LMO solves

$$s_t = \arg \min_{\|s\|_1 \leq R} \langle g_t, s \rangle.$$

By Hölder's inequality, the minimum is $\langle g_t, s \rangle \geq -\|g_t\|_\infty \cdot \|s\|_1 \geq -R \|g_t\|_\infty$, achieved by

$$s_t = -R \cdot \text{sign}(g_{t,j^*}) \cdot e_{j^*}, \quad j^* = \arg \max_i |g_{t,i}|.$$

This is an $\mathcal{O}(d)$ operation (one argmax pass) and returns a *single-nonzero* vector, ensuring every update direction is maximally sparse [6].

4.2 Variance-Controlled Step Size

The core contribution of FW-VC is an adaptive step size that responds to gradient variability. We maintain an EMA of the squared gradient norm,

$$v_t = \beta v_{t-1} + (1 - \beta) \|g_t\|^2,$$

and define the step size as

$$\gamma_t = \text{clip}\left(\frac{\alpha}{\sqrt{v_t} + \varepsilon}, \gamma_{\min}, \gamma_{\max}\right).$$

Intuitively, when stochastic gradients are large and noisy (high v_t), γ_t is small and the update is conservative; as training progresses and gradients stabilize (low v_t), γ_t grows, allowing the iterate to converge more aggressively onto a sparse vertex. The clipping to $[\gamma_{\min}, \gamma_{\max}]$ ensures both a lower bound (preventing stagnation) and an upper bound (preventing instability).

4.3 Full Update Rule

Algorithm 1 FW-VC: Frank-Wolfe with Variance-Controlled Step Size

Require: $x_0 = 0, \quad v_0 = 0, \quad$ parameters $R, \beta, \alpha, \gamma_{\min}, \gamma_{\max}, \varepsilon$

- 1: **for** $t = 1, 2, \dots, T$ **do**
- 2: Sample mini-batch $\mathcal{B}$; compute g_t
- 3: $v_t \leftarrow \beta v_{t-1} + (1 - \beta) \|g_t\|^2$
- 4: $\gamma_t \leftarrow \text{clip}(\alpha/(\sqrt{v_t} + \varepsilon), \gamma_{\min}, \gamma_{\max})$
- 5: $s_t \leftarrow -R \cdot \text{sign}(g_{t,j^*}) \cdot e_{j^*} \quad [j^* = \arg \max_i |g_{t,i}|]$
- 6: $x_{t+1} \leftarrow (1 - \gamma_t)x_t + \gamma_t s_t$
- 7: **end for**

We also implement **SFW-Paper**, the classical stochastic FW with the open-loop schedule $\gamma_t = 2/(t+2)$ [6], as a projection-free reference baseline.

5 Theoretical Analysis

Proposition 1 (Feasibility preservation). *Suppose $\|x_0\|_1 \leq R$ and s_t is an ℓ_1 -ball vertex with $\|s_t\|_1 = R$. Then for any $\gamma_t \in [0, 1]$, $x_{t+1} = (1 - \gamma_t)x_t + \gamma_t s_t$ satisfies $\|x_{t+1}\|_1 \leq R$.*

Proof. By convexity of the ℓ_1 norm [1]: $\|x_{t+1}\|_1 \leq (1 - \gamma_t)\|x_t\|_1 + \gamma_t\|s_t\|_1 \leq (1 - \gamma_t)R + \gamma_t R = R$. Since $\gamma_t \in [\gamma_{\min}, \gamma_{\max}] \subset [0, 1]$, this holds at every iteration. $\square$

Proposition 2 (Sparsity accumulation). *Starting from $x_0 = 0$, after T iterations of FW-VC, x_T is a convex combination of at most T vertices of the ℓ_1 ball [6]. Since each vertex has support size one, $\|x_T\|_0 \leq T$.*

In practice, sparsity is far below this bound: the LMO repeatedly selects the same dominant coordinate when the gradient direction is consistent, concentrating weight onto few nonzero entries.

Convergence. For smooth convex f , the standard stochastic FW with schedule $\gamma_t = \mathcal{O}(1/\sqrt{t})$ achieves $\mathcal{O}(1/\sqrt{T})$ expected suboptimality [5]. Our adaptive schedule satisfies the Robbins-Monro conditions [10]: $\sum_t \gamma_t \geq T \cdot \gamma_{\min} \rightarrow \infty$ (sufficient descent), while $\gamma_t \leq \gamma_{\max}$ (bounded step). The EMA v_t tracks gradient magnitude, acting as an online estimate of the noise level. When v_t is large (noisy early phase), γ_t is small, matching conservative early-phase behavior; as gradients stabilize, γ_t increases, analogous to

RMSProp [12] adapted for convex combination updates.

FW duality gap. We track the Frank-Wolfe gap $\delta_t = \langle g_t, x_t - s_t \rangle$ as a convergence diagnostic. A gap of zero certifies first-order stationarity; for convex f , $\delta_t \rightarrow 0$ implies $x_t \rightarrow x^*$ [6].

6 Experimental Setup

6.1 Datasets

We evaluate on two binary classification benchmarks, both from `scikit-learn` [8].

Digits (MNIST-proxy). The 8×8 handwritten digit dataset (64 features, 1,797 samples). We define a binary task: digit 0 versus all others ($\approx 10\%$ positive rate). This serves as an image classification proxy for the MNIST experiment described in our proposal.

Wisconsin Breast Cancer. A tabular dataset with 30 clinical measurements (569 samples). The binary task is malignant versus benign (37% positive rate). This replaces the UCI Adult dataset from the original proposal, which was unavailable due to network access constraints; the Breast Cancer dataset provides a comparable tabular benchmark for sparsity evaluation.

All features are standardized to zero mean and unit variance using training-set statistics, and a bias column is appended.

6.2 Baselines

We compare five methods in total. **SGD** (unconstrained mini-batch stochastic gradient descent) and **Adam** [7] serve as the unconstrained accuracy upper bounds. **Projected SGD** applies Euclidean projection onto the ℓ_1 ball after each gradient step, using the $\mathcal{O}(d \log d)$ isotonic algorithm [2]. **SFW-Paper** is the classical stochastic FW with open-loop schedule $\gamma_t = 2/(t+2)$ [6]; it serves as the projection-free reference. **FW-VC** is our proposed method.

6.3 Hyperparameters and Evaluation

All methods run for 40 epochs. Mini-batch sizes are 64 (Digits) and 32 (Breast Cancer). SGD and Projected SGD use $\eta = 0.05$; Adam uses $\eta = 0.01$ (tuned for fair convergence). The ℓ_1 radius R is shared

across constrained methods. FW-VC hyperparameters were selected by grid search; Table 1 summarizes the configuration.

Table 1: FW-VC hyperparameters per dataset.

Parameter	Digits	Breast Cancer
R (radius)	3.0	5.0
β (EMA decay)	0.90	0.95
α (scale)	1.0	2.0
$\gamma_{\max}$	0.10	0.10
$\gamma_{\min}$	0.001	0.001

We report training and test cross-entropy loss, training and test accuracy, sparsity ratio (fraction of $|w_i| \leq 10^{-6}$), number of nonzero weights (NNZ), ℓ_1 norm, and wall-clock runtime. We additionally track the FW duality gap for projection-free methods and the adaptive step size trajectory γ_t for FW-VC.

7 Results

7.1 Digits Dataset

Table 2 reports final performance on the Digits dataset. Among all methods, SGD and Adam achieve the highest accuracy (99.4% and 99.7% respectively) but use all 62 out of 65 weight coordinates. Of the three constrained methods, **FW-VC achieves the highest sparsity** (75.4%, using only 16 nonzero weights), with SFW-Paper producing 38 nonzero weights (41.5% sparsity) and Projected SGD achieving 17 (73.8% sparsity). Compared to Projected SGD — the natural constrained baseline — FW-VC has comparable sparsity at a lower per-step computational cost ($\mathcal{O}(d)$ LMO versus $\mathcal{O}(d \log d)$ projection). SFW-Paper achieves a better accuracy-sparsity balance (98.3% at 41.5% sparsity), reflecting that its open-loop schedule makes more conservative weight assignments early in training.

Table 2: Final results on the Digits dataset (40 epochs). Sparsity = fraction of $|w_i| \leq 10^{-6}$. Constrained methods (rows 3–5) enforce $\|x\|_1 \leq R$.

Method	Test Loss	Test Acc	Sparsity	NNZ
SGD	0.0478	99.4%	0.046	62
Adam	0.0409	99.7%	0.046	62
Proj. SGD	0.2015	90.8%	0.738	17
SFW-Paper	0.2162	98.3%	0.415	38
FW-VC	0.2201	93.1%	0.754	16

7.2 Breast Cancer Dataset

Table 3 shows results on the Breast Cancer benchmark. Here, unconstrained methods (SGD, Adam) achieve 98.2% accuracy using all 31 features. Among constrained methods, Projected SGD and SFW-Paper both reach 98.2% test accuracy, with SFW-Paper using nearly all features (30 NNZ). FW-VC achieves 93.9% accuracy with 22 nonzero weights (29.0% sparsity), the highest sparsity among all five methods. The larger accuracy gap here — relative to the Digits dataset — is attributable to the dataset’s structure: all 30 breast cancer features are clinically meaningful, so enforcing sparsity through ℓ_1 constraints unavoidably discards some predictive information. This is an expected and honest characteristic of constrained sparse methods on datasets with no redundant features.

Table 3: Final results on the Breast Cancer dataset (40 epochs).

Method	Test Loss	Test Acc	Sparsity	NNZ
SGD	0.0814	98.2%	0.000	31
Adam	0.0710	98.2%	0.000	31
Proj. SGD	0.1373	98.2%	0.194	25
SFW-Paper	0.1724	98.2%	0.032	30
FW-VC	0.2240	93.9%	0.290	22

7.3 Training Dynamics

Figure 1 shows the six-panel training dashboard for the Digits dataset. The step size panel (bottom right) confirms that γ_t for FW-VC is *not* trivially clamped to $\gamma_{\max}$ throughout training — it shows meaningful variation across epochs, demonstrating that the variance-control mechanism is active. The SFW-Paper step size decays monotonically following the $2/(t+2)$ schedule, reaching very small values (≈ 0.002) by epoch 40; this makes SFW-Paper increasingly conservative in late training, while FW-VC maintains a more aggressive step.

The sparsity panel shows FW-VC accumulating sparsity fastest among constrained methods and reaching the highest final sparsity ratio on both datasets.

7.4 Convergence Stability

Figure 3 shows the rolling standard deviation of training loss (window = 5 epochs). SGD and Adam exhibit the most stable convergence. Among projection-free methods, SFW-Paper is more stable




Figure 1: Training dashboard for the Digits dataset. Each panel shows one metric across all five methods. The bottom-right panel shows the per-epoch mean step size γ_t for the two FW methods.




Figure 2: Frank-Wolfe duality gap across training on the Digits dataset. FW-VC maintains a slightly higher gap than SFW-Paper, reflecting its larger step sizes in late training; both gaps decrease across epochs, consistent with convergence toward the constrained optimum.

than FW-VC, particularly in early training, because its step size decays monotonically while FW-VC’s adaptive rule may produce larger steps when v_t is small early on. FW-VC stabilizes by mid-training, and both FW methods have comparable late-phase stability.

7.5 Hyperparameter Sensitivity

Figure 4 shows a one-at-a-time sensitivity sweep for FW-VC on the Digits dataset. The method is most sensitive to α and $\gamma_{\max}$, which jointly control the effective step size. It is relatively robust to R (for $R \geq 2$) and β (for $\beta \in [0.85, 0.95]$). This indicates that practitioners can tune β and R with reasonable confidence, while α and $\gamma_{\max}$ benefit from targeted tuning.




Figure 3: Rolling standard deviation of training loss (window = 5) on Digits. Lower values indicate more stable convergence.




Figure 4: One-at-a-time hyperparameter sensitivity of FW-VC on Digits. Each subplot varies one parameter while holding others at their tuned values.

8 Discussion

FW-VC vs. SFW-Paper. The comparison between our proposed method and the classical stochastic FW is the most informative. On Digits, FW-VC achieves higher sparsity (75.4% vs. 41.5%) but lower accuracy (93.1% vs. 98.3%). This tradeoff arises because FW-VC’s larger adaptive step size drives iterates more aggressively toward sparse vertices, while SFW-Paper’s decaying open-loop schedule assigns weight more gradually across coordinates, retaining more of them. Neither approach dominates the other: FW-VC is preferable when sparsity is the primary objective, SFW-Paper when accuracy is more important in the constrained regime.

FW-VC vs. Projected SGD. Both enforce $\|x\|_1 \leq R$ but through different mechanisms. FW-VC achieves comparable or higher sparsity than Projected SGD on both datasets, at strictly lower per-step cost. The accuracy gap (93.1% vs. 90.8% on Digits; 93.9% vs. 98.2% on Breast Cancer) illustrates that neither method consistently dominates, but FW-VC’s advantage in sparsity and computational simplicity makes it the more principled choice for settings where the ℓ_1 constraint is tight.

Does the adaptive step size work? The γ_t trajectory plots confirm that the variance-control mech-

anism is genuinely active: the step size is not trivially clamped at $\gamma_{\max}$ throughout training but varies in response to gradient dynamics. This validates our key design choice and distinguishes FW-VC from a simpler constant-step FW.

Limitations. The method introduces three additional hyperparameters (α , β , $\gamma_{\max}$) beyond the ℓ_1 radius R . Sensitivity analysis shows this tuning effort is manageable, but it is a practical consideration. Additionally, the convergence rate guarantee for our specific adaptive schedule is informal — a rigorous rate characterization (analogous to AdaGrad’s $\mathcal{O}(\ln T/\sqrt{T})$ guarantee) would strengthen the theoretical contribution.

9 Implementation Notes

The implementation is organized into a clean multi-module Python project with separate packages for data loading (`src/data`), constrained optimization primitives and optimizers (`src/optimizers`), experiment orchestration and sensitivity analysis (`src/experiments`), and shared numerical utilities (`src/utills`). The entry point `run.py` reproduces all results and figures. All code is available at <https://github.com/VaishnaviNasery/fw-vc-project>.

Numerical stability is ensured throughout: logits are clipped to $[-50, 50]$ before sigmoid evaluation, gradients are norm-clipped to 5.0, and NaN/Inf values are sanitized after each step. All experiments use `numpy.random.seed(42)` for reproducibility.

10 Conclusion

We presented FW-VC, a variance-controlled stochastic Frank-Wolfe algorithm for ℓ_1 -constrained machine learning. Our contributions are: a principled adaptive step size rule designed for the convex combination FW update structure; formal feasibility and sparsity guarantees; and a comprehensive empirical evaluation across five optimizers and two datasets.

FW-VC produces the sparsest weight vectors among all methods, using only 16 of 65 weights (75.4% sparsity) on the Digits task while maintaining 93.1% test accuracy. It achieves this with $\mathcal{O}(d)$ per-step cost, strictly cheaper than the $\mathcal{O}(d \log d)$ Projected SGD baseline. The adaptive step size is demonstrably active across training, confirming that

variance control adds meaningful behavior beyond a fixed or open-loop schedule.

Future work includes extending FW-VC to multi-class and neural network settings [9], deriving formal convergence rates for the adaptive schedule, and evaluating on larger-scale sparse recovery tasks where the computational advantage of the LMO over projection is most pronounced.

References

- [1] S. Boyd and L. Vandenberghe. *Convex Optimization*. Cambridge University Press, 2004. §3.1.1 used for norm convexity in feasibility proof.
- [2] J. Duchi, S. Shalev-Shwartz, Y. Singer, and T. Chandra. Efficient projections onto the ℓ_1 -ball for learning in high dimensions. *Proceedings of the 25th ICML*, 272–279, 2008.
- [3] J. Duchi, E. Hazan, and Y. Singer. Adaptive subgradient methods for online learning and stochastic optimization. *JMLR*, 12:2121–2159, 2011.
- [4] M. Frank and P. Wolfe. An algorithm for quadratic programming. *Naval Research Logistics Quarterly*, 3(1–2):95–110, 1956.
- [5] E. Hazan and S. Kale. Projection-free online learning. *Proceedings of the 29th ICML*, 1843–1850, 2012.
- [6] M. Jaggi. Revisiting Frank-Wolfe: Projection-free sparse convex optimization. *Proceedings of the 30th ICML*, 427–435, 2013.
- [7] D. P. Kingma and J. Ba. Adam: A method for stochastic optimization. *ICLR*, 2015.
- [8] F. Pedregosa et al. Scikit-learn: Machine learning in Python. *JMLR*, 12:2825–2830, 2011.
- [9] S. Pokutta, C. Spiegel, and M. Zimmer. Deep neural network training with Frank-Wolfe. *arXiv:2010.07243*, 2020.
- [10] H. Robbins and S. Monro. A stochastic approximation method. *Annals of Mathematical Statistics*, 22(3):400–407, 1951.
- [11] R. Tibshirani. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society, Series B*, 58(1):267–288, 1996.
- [12] T. Tieleman and G. Hinton. Lecture 6.5 — RM-SPop. *COURSERA: Neural Networks for Machine Learning*, 2012.